

Build a Healf Product Intelligence Agent

Context

Healf is a UK-based health and wellness marketplace stocking 200+ brands across supplements, nutrition, and wellbeing. We're building an AI-native operating model where every team has AI agents as functional colleagues — each with their own capabilities, tools, and documentation.

The Task

Build a functional MVP of a natural language agent that can answer questions about product pages on healf.com. A user should be able to give the agent a product URL and ask questions about it in plain English — and the agent should respond with useful, accurate answers.

The goal is a conversational agent that works today and can be improved with more tooling and capabilities over time. We're not looking for a finished product — we're looking for a solid foundation and a clear view of where it goes next.

For example:

https://healf.com/en-uk/products/lmnt-recharge-electrolytes-variety-pack?variant=44024590893295&selling_plan=6835863791

Healf's storefront runs on Shopify. There is structured data available beyond what's visible on the page — how you discover and use that is part of the test.

Across the site, there are implicit standards — how the strongest listings are written, how brands present themselves, what good looks like in terms of imagery and information density. An agent that understands the broader context of the site will give better answers about any individual page.

Required Capabilities

Your agent must have at least 4 distinct capabilities. How you architect these is up to you — they could be separate tools, modules, sub-workflows, or functions.

1. Navigate

Given a Healf product URL, the agent can go to that page and access its contents.

2. Ingest

The agent can extract and structure information from a product page. A PDP has many facets — product text and descriptions, customer reviews, images, pricing (one-time and subscription), ingredients, and more. For this MVP, your agent should be able to ingest at least 2 of: product text, reviews, or images. It's worth exploring whether Shopify exposes this data in a more structured way than scraping the rendered page.

3. Evaluate

Based on what the user asks, the agent can assess the relevant parts of the product page — whether that’s the quality of the description, the presence of reviews, the completeness of ingredient information, SEO signals, or anything else the question calls for. The best evaluations are grounded in context, not generic checklists.

4. Act on findings

The agent produces a useful response — answering the question directly, surfacing issues, making recommendations, or producing improved content where appropriate.

You may build additional capabilities beyond these four. What you choose to include in the MVP vs. what you leave for the improvement plan is part of the test.

Example Prompts

Here are some examples of the kind of questions your agent may need to handle:

“Does this product have any reviews?” [product URL]

“What can I improve on this page?” [product URL]

“Does this product have Vitamin D in it?” [product URL]

These are deliberately different in scope — factual lookups, open-ended evaluations, and specific product data questions. Show us how your agent handles at least a couple of these.

Deliverables

- Working MVP — this could be code, n8n workflow JSON, or any combination. It should produce real output from live healf.com data and accept natural language input.
- Example outputs — showing the agent’s actual responses to a few of the example prompts above
- README — your architecture decisions, what the MVP can do today, and a clear plan for how you’d extend it — what capabilities you’d add next, what you’d change for production, and what the agent looks like in 3 months

Constraints

- Time: Up to 5 hours. We’re respectful of your time — this is about a functional MVP and a plan for what comes next, not a polished product.
- Platform: Use whatever you want — Python, TypeScript, n8n, Make, LangChain, Claude Code, or any combination. We care about the result, not the stack.
- Data: Everything must come from the live Healf website. No pre-provided datasets.
- LLM usage: At least one capability must use an LLM for reasoning or evaluation — not just data extraction.
- Interface: The agent must accept natural language questions. We should be able to give it a URL and ask something in plain English.

What We’re Looking For

A conversational agent that works — even if it only does a few things today — built in a way that makes it easy to add more. We're paying attention to how you scope an MVP, how your agent handles a question, and whether your improvement plan shows you understand what matters.

Submission

Submit a GitHub link, zip, or workflow export — whatever fits your approach. Include example outputs so we can review without re-running.

Questions? Email operations@opvia.in — but keep in mind that ambiguity is part of the test.