

Daran Menezes

+91 99 0205 8501 | daranmenezes91@gmail.com | [linkedin.com/in/daran-menezes-6b6a99220](https://www.linkedin.com/in/daran-menezes-6b6a99220) | github.com/Danez-AI |
Mangalore, India (IST, UTC+5:30 — ~5h UK daily overlap)

PROFESSIONAL SUMMARY

AI automation engineer who designs, builds, and owns production-grade n8n workflows and AI agent pipelines end-to-end. Four self-deployed systems in daily production use, including a Human-In-The-Loop invoice pipeline for a live pharmacy business (three-tier LLM matching cascade, 116 tests, ~45 operator-hours saved per month) and an internal AI knowledge layer at Epsilon (n8n + RAG pipeline ingesting 50+ SOPs, enabling self-serve ops answers across the team). Fluent in Anthropic Claude Code, LLM API integration, REST APIs, OAuth, and webhook architectures. Brings 5 years of enterprise marketing automation depth — the production instinct that knows what breaks and how to design fault-tolerant workflows that hold. Available global remote (IST, UTC+5:30, ~5h UK daily overlap). Portfolio: github.com/Danez-AI/personal-trainer-agent

CORE COMPETENCIES

n8n Production Workflows

Agentic Pipeline Design

Anthropic & LLM API Integration

Human-In-The-Loop (HITL)

REST APIs & Webhook Architectures

Production Reliability & RCA

Process Discovery & ROI Scoping

AI Knowledge Layer & Enablement

WORK EXPERIENCE

Epsilon 01/2023 – Present

Campaign Business Analyst

- Built Epsilon's first internal **AI knowledge layer** — n8n + OpenAI Embeddings + GPT-4o mini RAG pipeline ingesting 50+ campaign SOPs — enabling self-serve answers for the ops team and reducing manual SOP search time across daily campaign operations.
- Architected an enterprise-internal **n8n workflow automation** replacing Okta SAML SSO manual proof triggering in Epsilon PeopleCloud Messaging: Playwright session persistence within the authenticated browser context (the only viable architecture for SSO-gated APIs), Gemini Vision OCR for screenshot-based key extraction, Docker Compose deployment behind Cloudflare Tunnel, append-only JSONL audit log; task time from several minutes to ~10 seconds, full delegation to any team member without credentials.
- Designed automated segmentation, content assembly, and nurture logic across email, push, and SMS; eliminated manual intervention at scale and reduced campaign technical debt across global lifecycle programs.
- Led end-to-end execution of multichannel campaigns for Fortune 500 clients (Marriott, AbbVie, McKinsey, National Geographic) — A/B testing, audience segmentation, performance optimization, and SLA compliance across time zones.
- Mentored junior analysts, drove recruitment, and served as primary escalation contact for client process improvements and platform issues.

Epsilon 08/2021 – 12/2022

Campaign Analyst

- Executed high-volume campaign builds and deployments across email and push; troubleshooted platform issues for rapid resolution.
- Managed HTML/CSS segmentation and personalization logic within Braze and EPCM for retail, pharma, electronics, and banking clients.
- Coordinated intake and project management via Jira and ServiceNow; supported data-driven reporting and campaign analysis.

TECHNICAL PROJECTS

SoftLife Automation — HITL Invoice Pipeline PRODUCTION

Designed and deployed a production n8n + Python automation replacing fully manual daily invoice processing for a local pharmacy business (5–15 invoices/day, 15–20 line items each, ~45 operator-hours saved per month). Architected a three-tier LLM matching cascade: SQLite correction cache (microsecond lookups for repeat items) → rapidfuzz token-sort against ~21,000 master items → Google Gemini for ambiguous cases. Implemented a **Human-In-The-Loop (HITL) gate** in n8n — operator reviews uncertain matches in a browser form before any import CSV is generated; enforces the invariant that all new items are confirmed before the batch runs (system rejects on violation). Automated entry into SoftLife's legacy PowerBuilder UI via pywinauto across cross-subnet machines (Tailscale + SSH). 116 tests passing. In daily production use.
Python 3.11 · n8n (self-hosted) · Google Gemini · rapidfuzz · pywinauto · SQLite · Pydantic v2 · Tailscale · Docker

Epsilon PCM Proof Automation — Enterprise Internal Tool PRODUCTION

Architected an end-to-end **n8n workflow automation** replacing manual proof email triggering in Epsilon PeopleCloud Messaging. Core constraint: Okta SAML SSO binds the session to the browser — extracted cookies are rejected externally. Solution: keep a live Playwright browser context running persistently and issue all PCM REST API calls from within the authenticated session (the only viable architecture). Built a 3-page n8n trigger workflow (BU selection → campaign preview confirmation → fire-and-forget with live status link), Gemini Vision OCR for key extraction from screenshots, and a webhook-driven auto-polling status page. Deployed as two Docker containers behind Cloudflare Tunnel with an append-only JSONL audit log for every proof trigger. Task time: several minutes → ~10 seconds; full delegation to any team member without credentials.

Node.js · Express · Playwright · n8n · Google Gemini 2.5 Flash · Docker Compose · Cloudflare Tunnel

GodFather — Self-Hosted AI Fitness Coach Agent PRODUCTION

Architected a zero-cost **agentic workflow** operating entirely through Telegram on a self-hosted home server. Four food-logging input modes (text, meal photo via Gemini Vision, barcode scan, nutrition label OCR) all routed through n8n. Five scheduled cron workflows (morning motivation, workout reminder, two meal reminders, weekly progress summary). Supabase (PostgreSQL + pgvector) stores food logs, weight history, and user profiles with structured queries for daily totals and trends. All services in Docker Compose behind Cloudflare Tunnel. Public repo:

github.com/Danez-AI/personal-trainer-agent

n8n · Supabase (PostgreSQL + pgvector) · Telegram Bot API · Google Gemini 2.0 Flash · Docker Compose · Cloudflare Tunnel

CCTV AI Insights — Cost-Aware Two-Stage Detection Pipeline PRODUCTION

Designed a two-stage AI pipeline for live CCTV: YOLOv8n runs continuous on-device frame inference at near-zero cost — Gemini Vision called only on confirmed detections for natural-language scene descriptions, preventing inference cost explosion. Debounce logic (one alert per class per 30s window), n8n-orchestrated Telegram delivery. Containerised as a hardened Docker service (non-root uid 1000, all capabilities dropped, read-only root filesystem).

Python · YOLOv8n · OpenCV · Google Gemini Vision · n8n · Docker Compose

EDUCATION

Bachelor of Science — St. Aloysius College, Mangalore

2021

SKILLS

Automation & Orchestration: n8n (self-hosted, Docker), webhook architectures, cron-driven agents, Human-In-The-Loop (HITL) design, Playwright

AI & LLM: Anthropic Claude Code (daily use), Google Gemini (Vision + reasoning), OpenAI Embeddings + GPT-4o, RAG pipelines, agentic workflows, YOLOv8

APIs & Integration: REST APIs, OAuth, JSON, webhooks, MCP (Model Context Protocol), Shopify, Slack, Airtable, Notion

Languages: JavaScript/Node.js, Python — working knowledge, AI-assisted development

Data & Storage: PostgreSQL + pgvector, Supabase, SQLite, Pydantic v2, fuzzy matching (rapidfuzz)

Infrastructure: Docker Compose, Cloudflare Tunnel, Tailscale VPN, SSH (Ed25519), Ubuntu 24.04 LTS

MarTech: Braze, EPCM, Shopify, Jira, ServiceNow